AWS - S3

What is Amazon S3?

- Amazon S3 stands for Simple Storage Service.
- It is cloud storage service provided by AWS
- S3 has characteristics -
  - Highly Scalable
  - Highly Available
  - Secure
  - Cost Effective
  - Performance
- It allows you to store and retrieve any amount of data from anywhere on the web.
- AWS S3 is also called as a version solution

What can you store in these S3?

- S3 services allows you to create buckets in which you can store everything.
  - Eg: we can store pictures, movies (or) any kind of files and folders (or) excel sheets etc..,
- But, as a DevOps Engineer What we will deal with is
  - Application log files, big databases, backup's, huge application supported configuration files etc..,
- Customers of all sizes and industries can use Amazon S3 to store and protect any amount of data for a range of use cases, such as data lakes, websites, mobile applications, backup and restore, archive, enterprise applications, IOT devices and big data analytics

What are S3 buckets?

S3 buckets are containers for storing objects (files) in Amazon S3. Each bucket has a unique name globally across all of AWS. You can think of an S3 bucket as a top-level folder that holds your data.

Buckets :

- Bucket is a component in S3
- If you create a folder in root level then it is called as buckets
- Buckets is used to store objects
- Bucket ownership is not transferrable to another account
- After we create a bucket, we can't change it's name (or) region
- We cannot delete the bucket directly. We need to make empty if first
- After a bucket is deleted, the name becomes available for reuse
- By default, we can create up to 100 buckets in each of your AWS account
- If you need additional buckets, you can increase your account bucket limit to 1,000 buckets by submitting a service limit increase

Objects :

- In bucket, you are storing the files is called objects
- To store your data in Amazon S3, you work with resources known as buckets and objects
- A bucket is a container for objects

- An object is a file and any metadata that describes that file
- With Amazon S3, you pay only for what you use.

Depending on the size of the data you are uploading, Amazon S3 offers the following options

- Upload an object with single operation using the AWS SDKs, REST API (or) AWS CLI
  - with single PUT operation, you can upload a single object up to 5GB in size
- Upload a single object using the Amazon S3 Console
  - With the Amazon S3 console, you can upload a single object up to 160 GB in size

Major Components Of Object :

- Key → name of the object. Eg: file.txt → key
- Value → Data in bytes. (file size)
- Version → Shows versioning ID for uniqueness of object
- Metadata → data about the data (or) it describes about the data

Why use S3 buckets?

- S3 buckets provide a reliable and highly scalable storage solution for various use cases.
- They are commonly used for backup and restore, data archiving, content storage for websites, and as a data source for big data analytics.
- It solves the storage problem for companies (or) organizations
- We can store the WAR files here
- Applications related log files generating purpose we can use S3
- Using S3 bucket we can host the website instead of EC2 instance
- Application contains some files. We can store that files here

We're using S3 instead of Nexus also, because maintenance easy in S3. and in nexus we need to do complex setup for that we need t2.medium, 20 GB volume, etc.., cost is high

S3 is globally accessible, we can access the data in s3 anywhere

What happened if S3 goes down ?

The Secret behind the success in S3 is → 11 (9's)

- 99.99999999999
  - It indicates the amount of reliability of AWS S3.
- i.e. when you upload a data in S3 and we thought like what if the data deleted/S3 down
  - So, S3 is giving the assurance like it giving the reliability to your objects
- So, that means in S3 our object will never gets deleted
  - Because of the replication mechanism of AWS
  - i.e. if you create an object S3 it will store your data in multiple availability zones. Every zone has some copies of your data. So, when you lose your data also, AWS have replicas. So, they can retrieve the data

Key Benefits of S3 buckets

S3 buckets offer several advantages, including:

- Durability and availability:
  - S3 provides high durability and availability for your data.
- Scalability:
  - You can store and retrieve any amount of data without worrying about capacity constraints.
- Security:
  - S3 offers multiple security features such as encryption, access control, and audit logging and bucket policies.
  - Encrypt data at rest using server-side encryption options provided by S3. Additionally, enable encryption in transit by using SSL/TLS for data transfers
  - Enable access logging to capture detailed records of requests made to your S3 bucket
  - Monitor access logs and configure alerts to detect any suspicious activities (or) unauthorized access attempts
- Performance:
  - S3 is designed to deliver high performance for data retrieval and storage operations.
  - If we create a S3 in the region that is near by us. then we can quickly access the content in S3. like you can upload (or) download the content in S3 very quickly
- Cost-effective:
  - S3 offers cost-effective storage options and pricing models based on your usage patterns.
  - It depends on the storage class patterns.
  - It depends on the type of data that you are storing and it depends on what your organization needs.
  - In the below image features are nothing but storage classes

Feature	S3 Standard	S3 Standard-IA	One Zone-IA	S3 Glacier	S3 Glacier Instant Retrieval	S3 Glacier Flexible Retrieval	S3 Glacier Deep Archive	S3 Outposts	S3 Intelligent-Tiering
Cost per GB per month	$0.02	$0.01	$0.01	$0.00	$0.00	$0.00	$0.00	$0.03	$0.015–0.025
Access time	1-15 seconds	3-5 minutes	3-5 minutes	12-48 hours	1-5 minutes	1-5 minutes	12-48 hours	Varies	Varies
Durability	100.00%	100.00%	100.00%	100.00%	100.00%	100.00%	100.00%	100.00%	100.00%
Availability	99.90%	99.90%	99.90%	99.90%	99.90%	99.90%	99.90%	99.90%	99.90%
Minimum storage duration	Varies	Varies	Varies	Varies	Varies	Varies	Varies	Varies	Varies

  - Eg: if we store 1TB data in S3 per one year. It will costs like 5 (or) 6 dollars.

Bucket properties and configurations

Versioning:

- Versioning allows you to keep multiple versions of an object in the bucket.
- It helps protect against accidental deletions or overwrites.

Bucket-level permissions and policies

Bucket-level permissions and policies define who can access and perform actions on the bucket. You can grant permissions using IAM (Identity and Access Management) policies, which allow fine-grained control over user access to the bucket and its objects.

Uploading and Managing Objects in S3 Buckets

Uploading objects to S3 buckets

You can upload objects to an S3 bucket using various methods, including the AWS Management Console, AWS CLI, SDKs, and direct HTTP uploads. Each object is assigned a unique key (name) within the bucket to retrieve it later.

Object metadata and properties

Object metadata contains additional information about each object in an S3 bucket. It includes attributes like content type, cache control, encryption settings, and custom metadata. These properties help in managing and organizing objects within the bucket.

File formats and object encryption

S3 supports various file formats, including text files, images, videos, and more. You can encrypt objects stored in S3 using server-side encryption (SSE). SSE options include SSE-S3 (Amazon-managed keys), SSE-KMS (AWS Key Management Service), and SSE-C (customer-provided keys).

Lifecycle management

Lifecycle management allows you to define rules for transitioning objects between different storage classes or deleting them automatically based on predefined criteria. For example, you can move infrequently accessed data to a lower-cost storage class after a specified time or delete objects after a certain retention period.

Multipart uploads

Multipart uploads provide a mechanism for uploading large objects in parts, which improves performance and resiliency. You can upload each part in parallel and then combine them to create the complete object. Multipart uploads also enable resumable uploads in case of failures.

Managing large datasets with S3 Batch Operations

S3 Batch Operations is a feature that allows you to perform bulk operations on large numbers of objects in an S3 bucket. It provides an efficient way to automate tasks such as copying objects, tagging, and restoring archived data.

Advanced S3 Bucket Features

S3 Storage Classes

S3 offers multiple storage classes, each designed for different use cases and performance requirements:

S3 Replication

S3 replication enables automatic and asynchronous replication of objects between S3 buckets in different regions or within the same region. Cross-Region Replication (CRR) provides disaster recovery and compliance benefits, while Same-Region Replication (SRR) can be used for data resilience and low-latency access.

S3 Event Notifications and Triggers

S3 event notifications allow you to configure actions when specific events occur in an S3 bucket. For example, you can trigger AWS Lambda functions, send messages to Amazon Simple Queue Service (SQS), or invoke other services using Amazon SNS when an object is created or deleted.

S3 Batch Operations

S3 Batch Operations allow you to perform large-scale batch operations on objects, such as copying, tagging, or deleting, across multiple buckets. It simplifies managing large datasets and automates tasks that would otherwise be time-consuming.

Security and Compliance in S3 Buckets

- S3 bucket security considerations
  - Ensure that S3 bucket policies, access control, and encryption settings are appropriately configured.
  - Regularly monitor and audit access logs for unauthorized activities.
- Data encryption at rest and in transit
  - Encrypt data at rest using server-side encryption options provided by S3.
  - Additionally, enable encryption in transit by using SSL/TLS for data transfers.
- Access logging and monitoring
  - Enable access logging to capture detailed records of requests made to your S3 bucket.
  - Monitor access logs and configure alerts to detect any suspicious activities or unauthorized access attempts.

S3 Bucket Management and Administration

- S3 bucket policies
  - Create and manage bucket policies to control access to your S3 buckets.
  - Bucket policies are written in JSON and define permissions for various actions and resources.
- S3 access control and IAM roles
  - Use IAM roles and policies to manage access to S3 buckets.
  - IAM roles provide temporary credentials and fine-grained access control to AWS resources.
- S3 APIs and SDKs
  - Interact with S3 programmatically using AWS SDKs or APIs.
  - These provide libraries and methods for performing various operations on S3 buckets and objects.
- Monitoring and logging with CloudWatch
  - Utilize Amazon CloudWatch to monitor S3 metrics, set up alarms for specific events, and collect and analyze logs for troubleshooting and performance optimization.
- S3 management tools
  - AWS provides multiple management tools, such as the AWS Management Console, AWS CLI, and third-party tools, to manage S3 buckets efficiently and perform operations like uploads, downloads, and bucket configurations.

Troubleshooting and Error Handling

- Common S3 error messages and their resolutions
  - Understand common S3 error messages like access denied, bucket not found, and exceeded bucket quota.

    - Troubleshoot and resolve these errors by checking permissions, bucket configurations, and network connectivity.
- Debugging S3 bucket access issues
    - Investigate and resolve issues related to access permissions, IAM roles, and bucket policies.
    - Use tools like AWS CloudTrail and S3 access logs to identify and troubleshoot access problems.
- Data consistency and durability considerations
    - Ensure data consistency and durability by understanding S3's data replication and storage mechanisms.
    - Verify that data is correctly uploaded, retrieve objects using proper methods, and address any data integrity issues.
- Recovering deleted objects
    - If an object is accidentally deleted, you can often recover it using versioning or S3 event notifications.
    - Additionally, consider enabling Cross-Region Replication (CRR) for disaster recovery scenarios.

Creating and Configuring S3 Buckets

Practice - DeepDive

Go to AWS → search S3 service → open → you get dashboard → click on create bucket

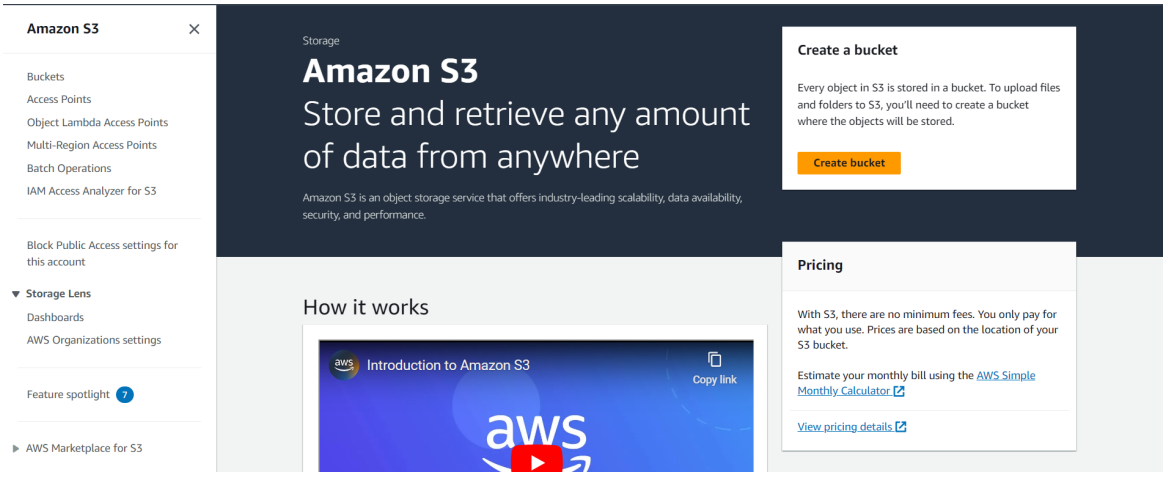

Creating an S3 bucket

- To create an S3 bucket, you can use the AWS Management Console, AWS CLI (Command Line Interface), or AWS SDKs (Software Development Kits).
- You need to specify a globally unique bucket name because it is globally accessing and select the region where you want to create the bucket.

Choosing a bucket name and region

- The bucket name must be unique across all existing bucket names in Amazon S3.
- It should follow DNS naming conventions, be 3-63 characters long, and contain only lowercase letters, numbers, periods, and hyphens.
- The region selection affects data latency and compliance with specific regulations.
- We are having regions, here we are choosing the nearest region due to latency issues
    - when I send a request to AWS, it has to cross multiple routers and it will reach data center and again I will get the request from AWS.

- So, usually if my region is Mumbai means, it's very near for me.
- To Solve the latency issues, we have to use region
- Any resource that you are creating on AWS are bound to region

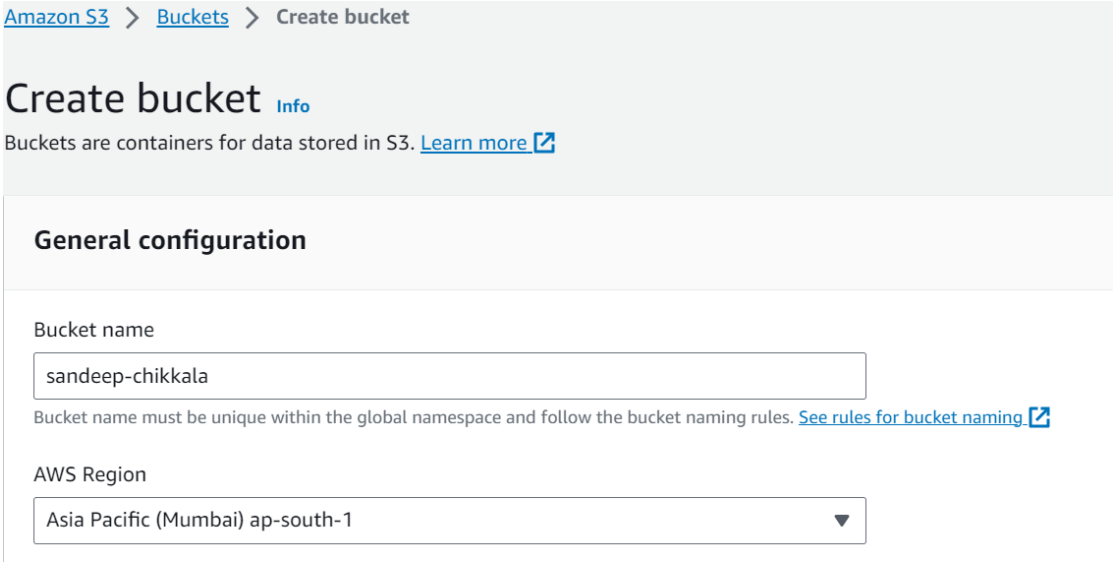

Object Ownership

- Objects can be accessed by only you means, select ACL disabled
- Objects can be accessible to all means, select ACL enabled

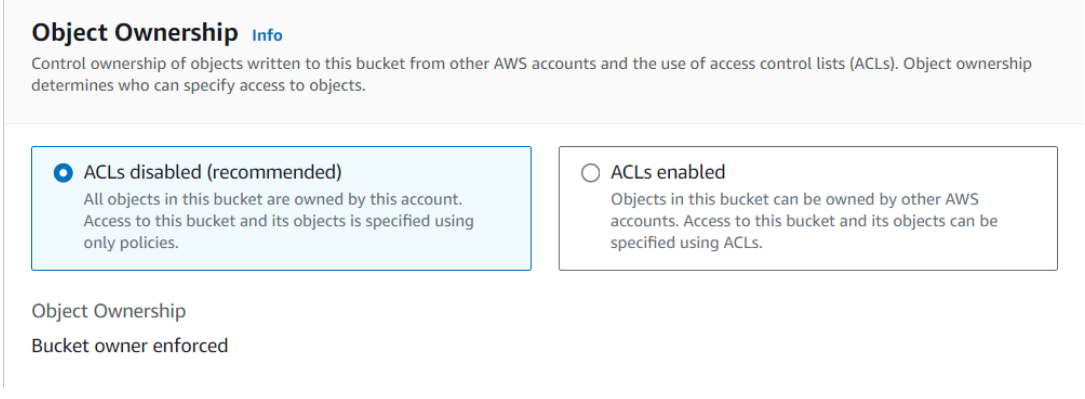

Block Public access

- Inside bucket objects will be access to outside, uncheck the block all public access

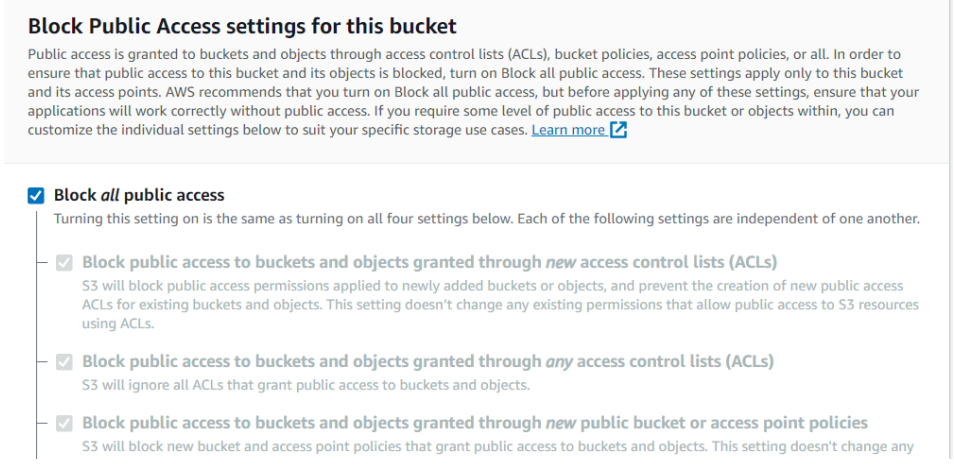

existing policies that allow public access to S3 resources.

- [x] **Block public and cross-account access to buckets and objects through *any* public bucket or access point policies**
  S3 will ignore public and cross-account access for buckets or access points with policies that grant public access to buckets and objects.

Bucket Versioning

Versioning is a means of keeping multiple variants of an object in the same bucket. You can use versioning to preserve, retrieve, and restore every version of every object stored in your Amazon S3 bucket. With versioning, you can easily recover from both unintended user actions and application failures. Learn more

Bucket Versioning

- (•) Disable
- ( ) Enable

Tags - *optional* (0)

You can use bucket tags to track storage costs and organize buckets. Learn more

No tags associated with this bucket.

Add tag

Default encryption Info

Server-side encryption is automatically applied to new objects stored in this bucket.

Encryption type Info

- (•) Server-side encryption with Amazon S3 managed keys (SSE-S3)
- ( ) Server-side encryption with AWS Key Management Service keys (SSE-KMS)
- ( ) Dual-layer server-side encryption with AWS Key Management Service keys (DSSE-KMS)
  Secure your objects with two separate layers of encryption. For details on pricing, see **DSSE-KMS pricing** on the **Storage** tab of the Amazon S3 pricing page.

Bucket Key

Using an S3 Bucket Key for SSE-KMS reduces encryption costs by lowering calls to AWS KMS. S3 Bucket Keys aren't supported for DSSE-KMS. Learn more

- ( ) Disable
- (•) Enable

▸ Advanced settings

After creating the bucket, you can upload files and folders to the bucket, and configure additional bucket settings.

Cancel | Create bucket

Once, click on Create bucket, you will get the below dashboard

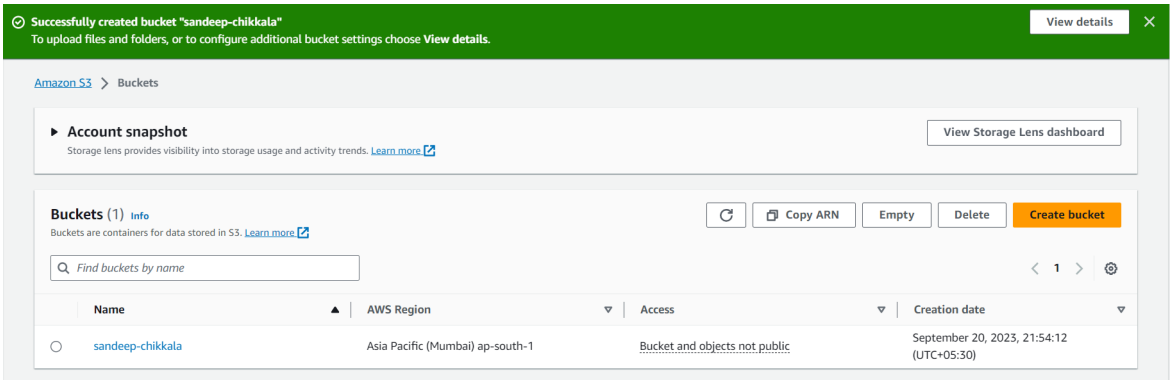

- Click on the bucket name, you will get below image

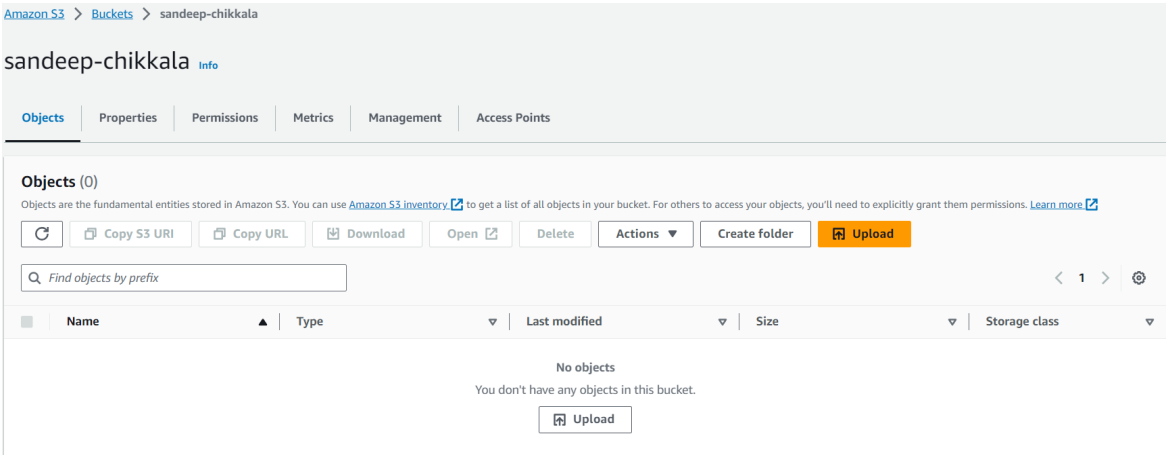

  - when you click on the bucket, usually you can see objects. Present we don't have objects
  - so, click on upload, you can click on add files and upload the data from your local
  - At a time, multiple files & folders also we can upload.

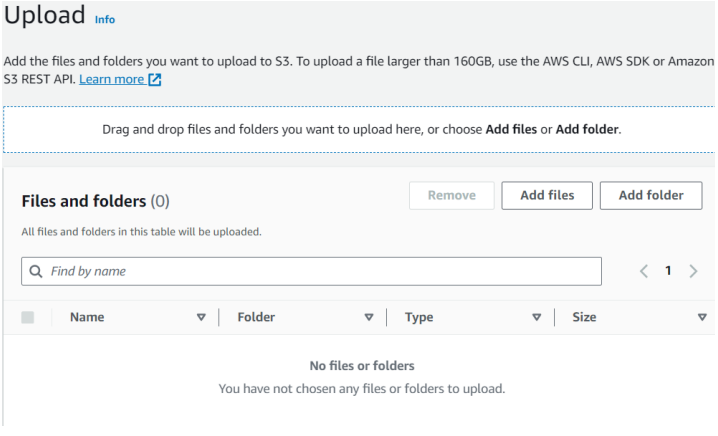

- After successfully uploaded, you will get below image

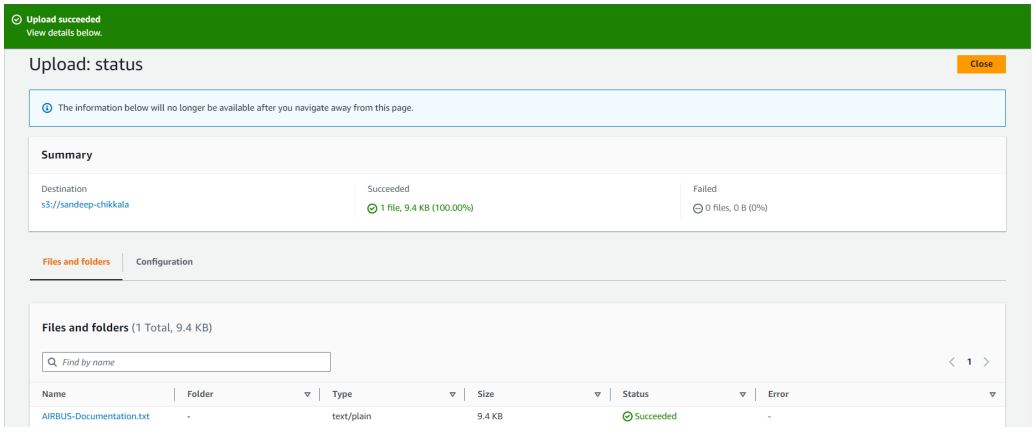

- After uploaded the files in the bucket, you will have the objects. If you select that fileg, you will get below image

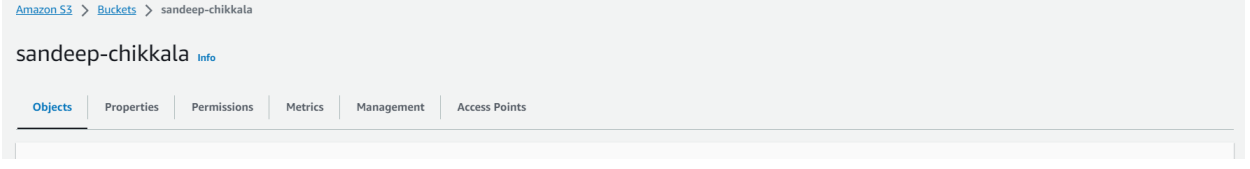

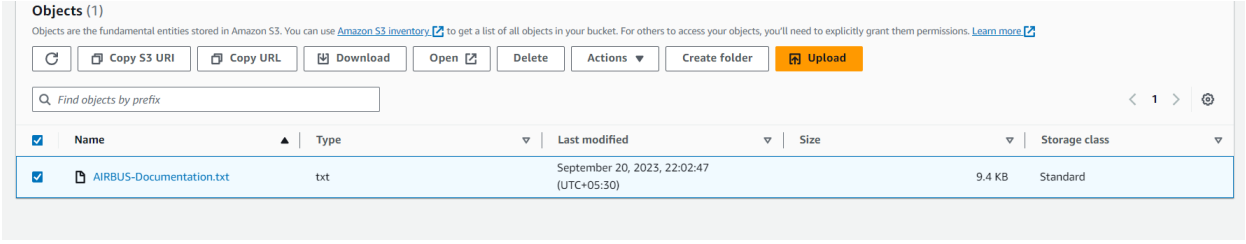

- If you select a file we're getting so many options
- Copy S3 URI → it contains file path → s3://sandeep-chikkala/AIRBUS-Documentation.txt
- Copy URL → It contains Bucket related URL
  - https://sandeep-chikkala.s3.ap-south-1.amazonaws.com/AIRBUS-Documentation.txt
  - when you access this URL in browser, you will get permission denied. Because, security is high for this

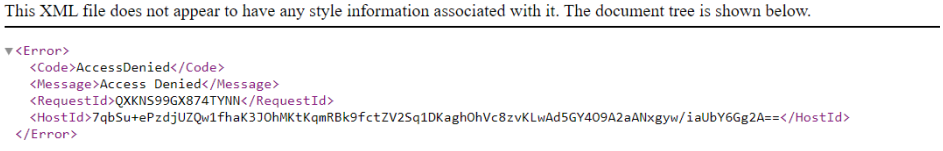

- Open → when you click on open, you will see the file output in browser
- Delete → we can delete the object file

If you want to delete a entire bucket, First, we need to empty the bucket. i.e. inside objects/files we have to delete, then after delete the bucket

- [ ] When you opened a bucket, we're having properties

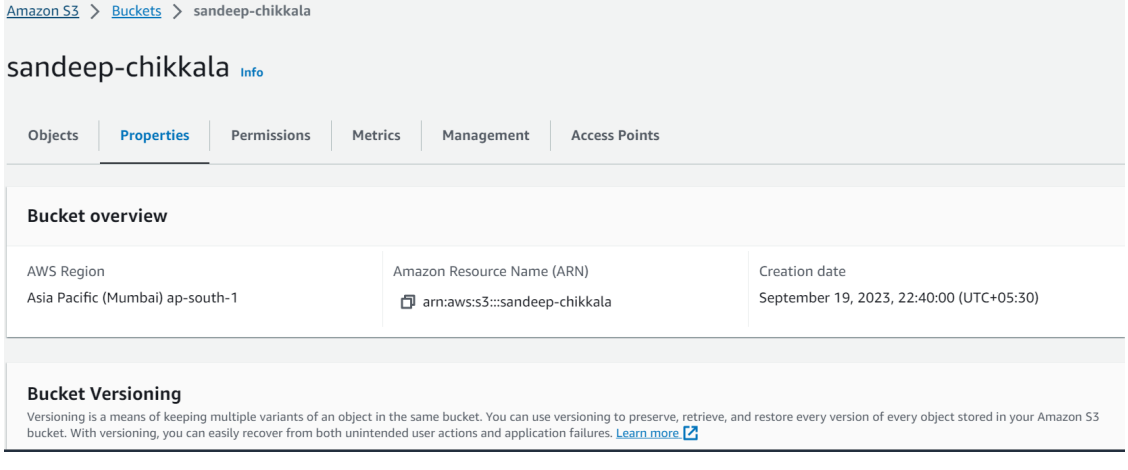

- If you want to enable the version, click on edit

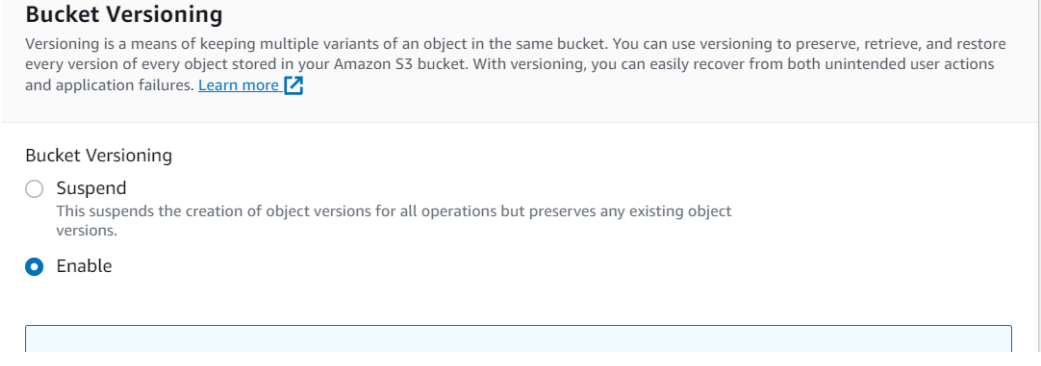

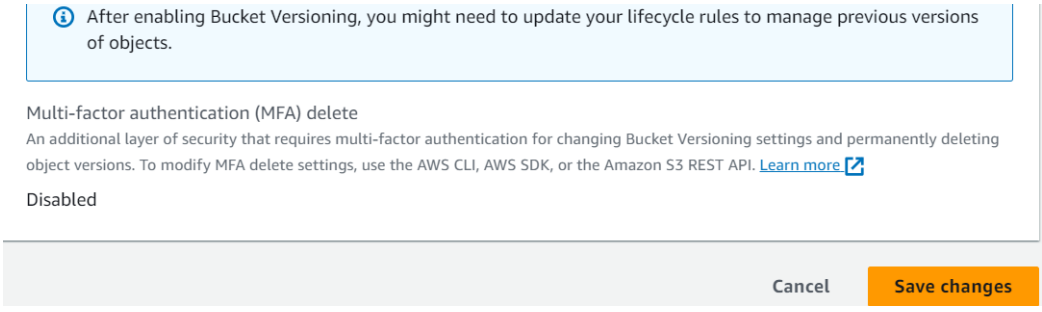

- Tags is used for identifying the resources

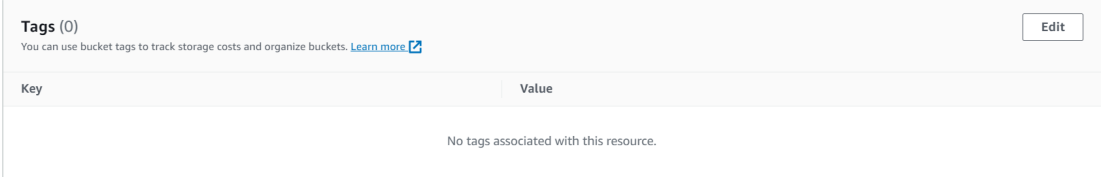

  - click on edit, and add the data like below image

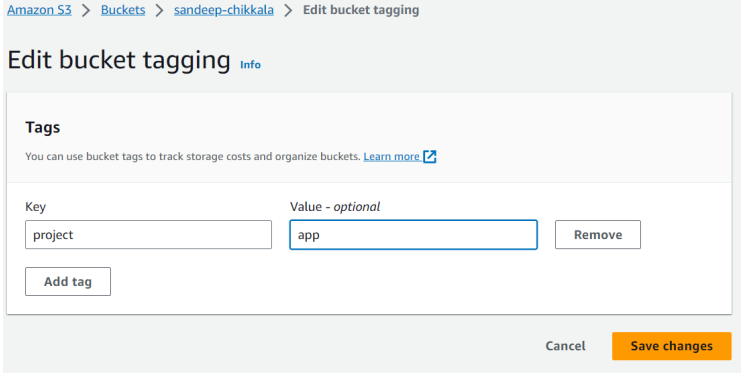

  - when exactly, we will use means let's say we having multiple ec2, buckets across the world. If someone says that hey can you give me a list of S3 buckets that is used by this specific project then we can make use of this tags to identifying purpose
  - Every organization or every project that you are working with they have to create tags for sure and that way it will be easy for us to extract

Default Encryption

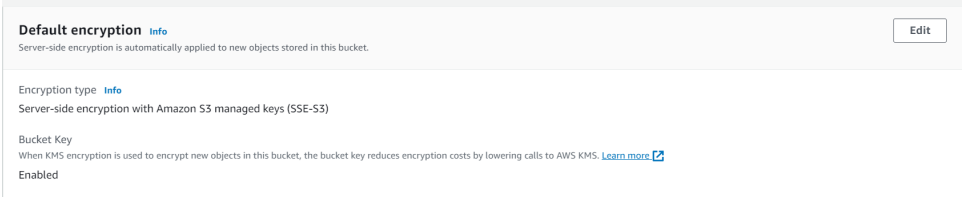

  - S3 recently added the default encryption. Previously, you should have enabled the encryption on the objects but now the bucket encryption is available by default

Server Access Logging

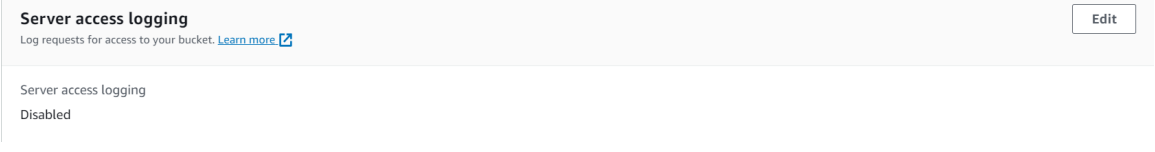

- We can enable this access logging by default this is disabled
- If you enable the access logging then it will ask like should be the policy here and then you can enable this access logging to ensure, who is logging into this bucket uh what kind of actions they are performing

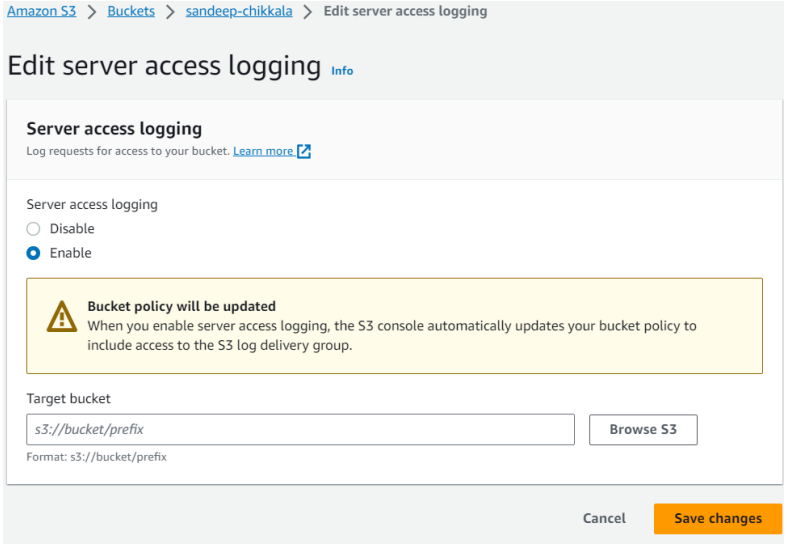

- we can restrict also
- We can remove their access and we can send out the notifications depending upon the actions that they are performing

Object Locking

- once we upload the object into the bucket. I want to lock the object so that nobody else should use that object and that is locked there should not be any updates to the object right
- If someone wants to update the object then they cannot update the object probably some sensitive information which you have decided that nobody has to override then you can use the bucket locking

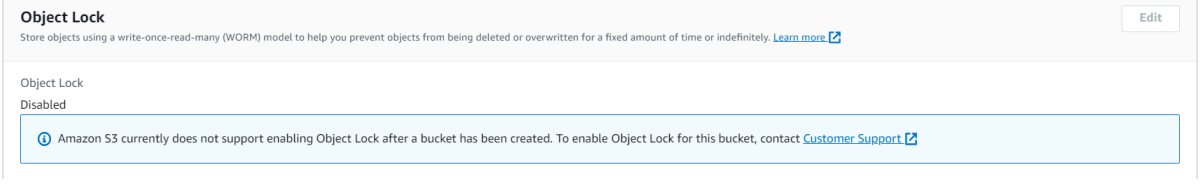

Enable the Version

How to enable the version for created bucket ?

- click on bucket → properties

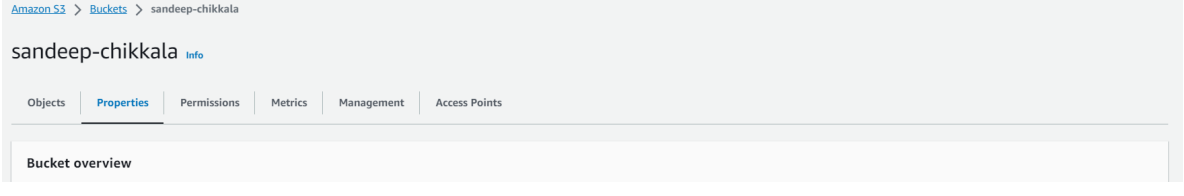

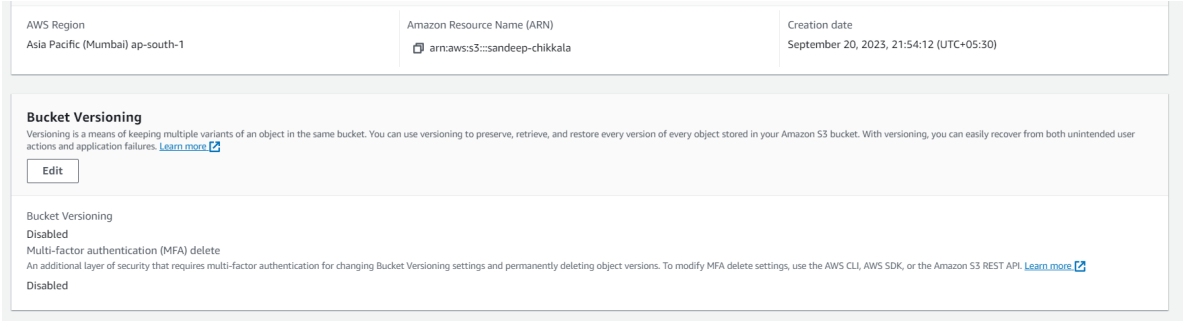

- **Now, go to Bucket Versioning → click on edit → click on enable → save**

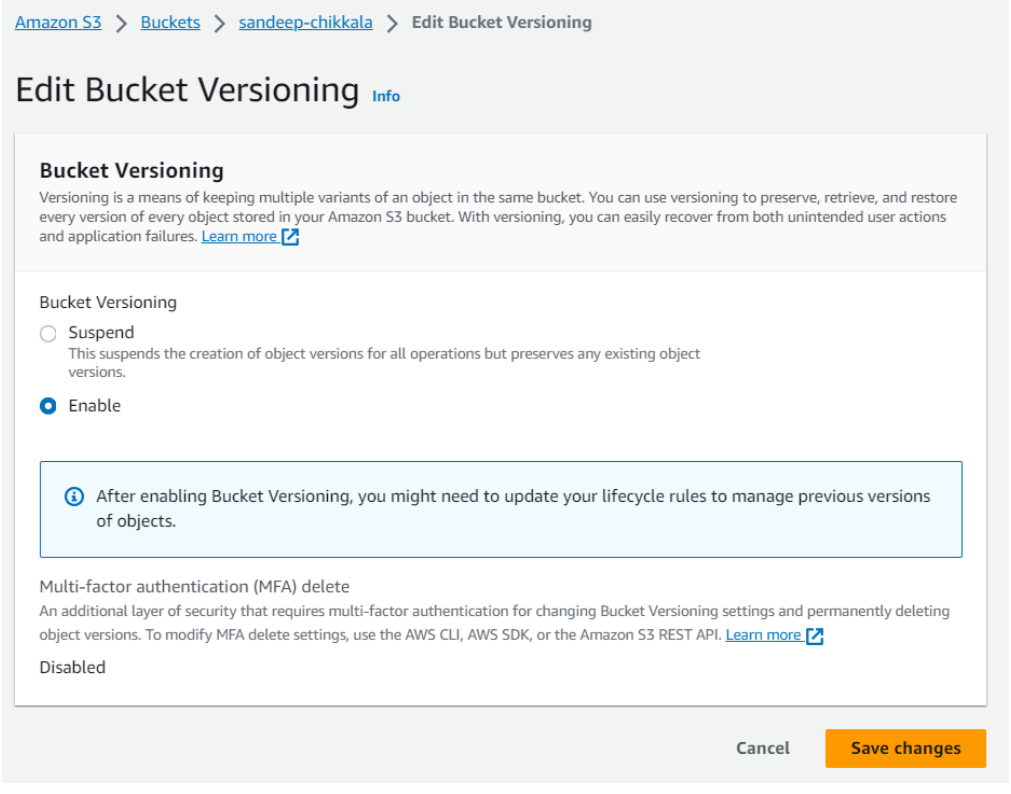

Now, when you go to buckets, you can see the "Show versions" in below image

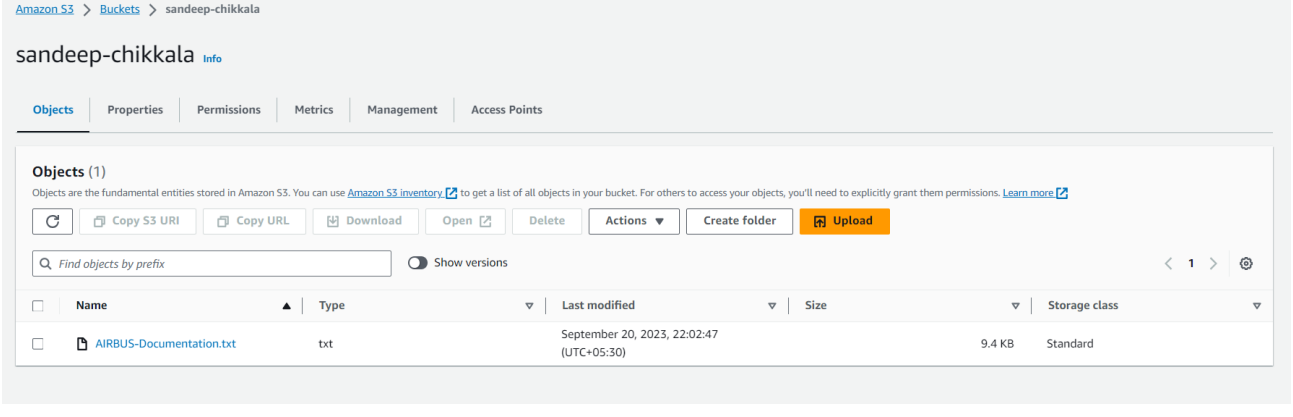

- **Now, when you delete the file means, it will deleted. But, when you enable the show versions, you can get the files like below image**

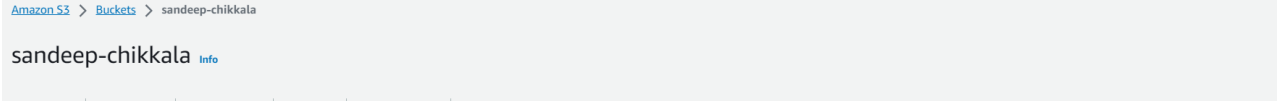

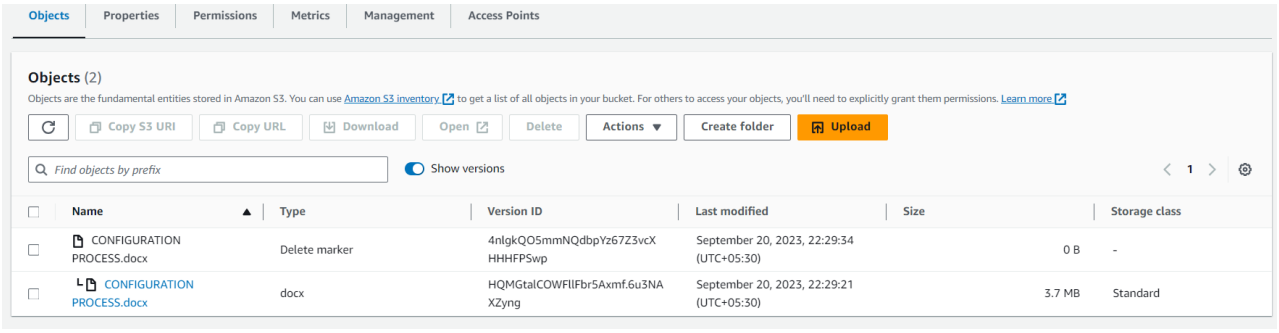

- Now, when you click on the "standard" file in storage class i.e. 2nd one, you will get below image

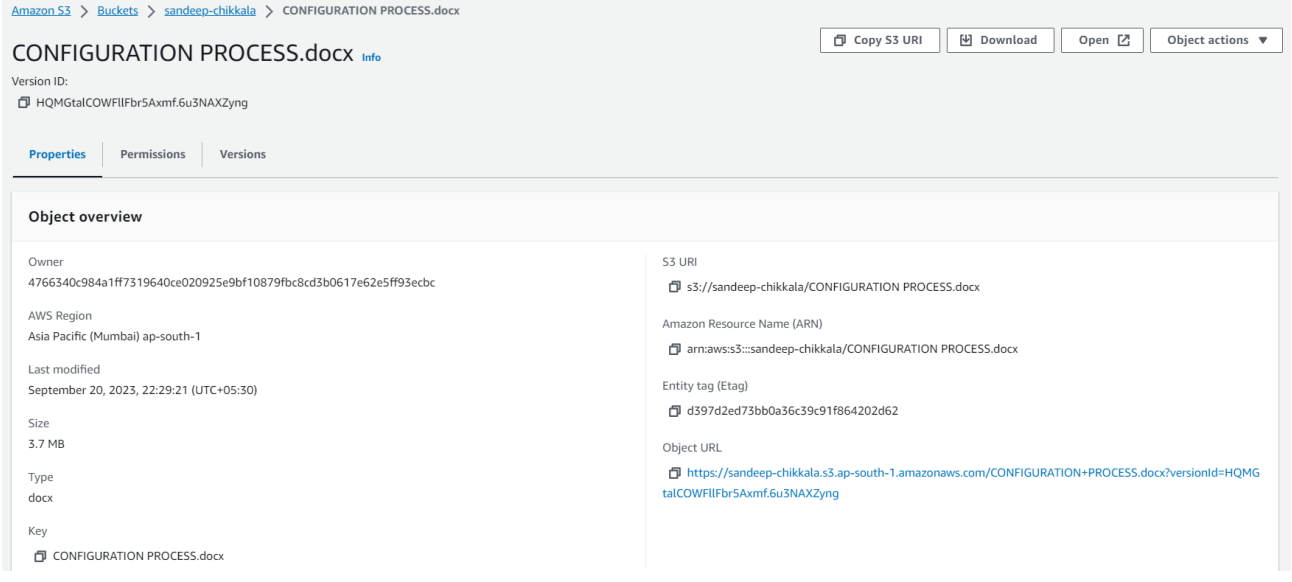

  - here, if you click on download, file will downloaded. And you can upload the file again

And when you're enable the versions, It will keep the multiple versions of the single file

- For ex: if you upload the same file, multiple times it will track every file

Enable the Object Ownership

- Here, we are enabling the ownership for created bucket

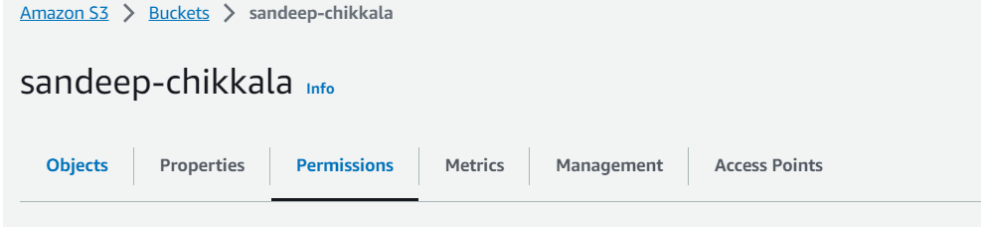

- Now, go to permissions → Go to ACL

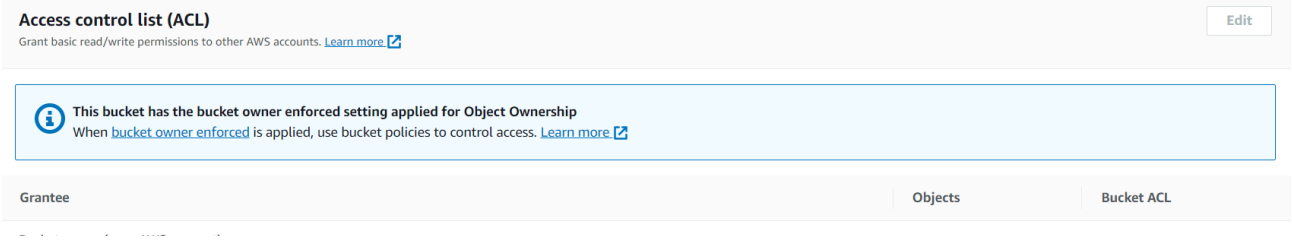

Bucket owner (your AWS account) Canonical ID: 4766340c984a1ff7319640ce020925e9bf10879fbc8cd3b0617e62e5ff93ecbc	List, Write	Read, Write
Everyone (public access) Group: http://acs.amazonaws.com/groups/global/AllUsers	-	-
Authenticated users group (anyone with an AWS account) Group: http://acs.amazonaws.com/groups/global/AuthenticatedUsers	-	-
S3 log delivery group Group: http://acs.amazonaws.com/groups/s3/LogDelivery	-	-

- Here, click on "bucket owner enforced"

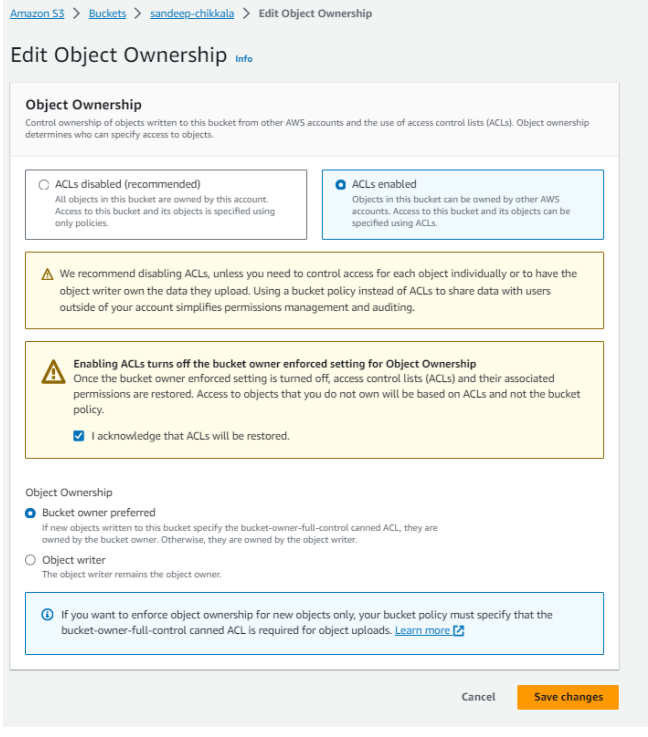

Advantages of Ownership

- If you enabled the ACL, previously when you're accessing the Copy URL from the object. It is showing permission denied error
- But, you enabled the ACL means, it will be open to all. So, people can read the file

CRR (Cross Region Replication)

If we upload the file in one bucket, that file will goes to another bucket,

Step - 1:

- Create 2 buckets, the names are
  - sandeep-sandy
  - sandeep-chikkala
- Enable the versioning, ACL enabled, Block all public access

Step - 2:

- Here, we have to enable the versioning for 2 buckets

- Go to 1st bucket → management → replication rule → create replication rule → replication name → status - enabled → Click on apply all objects in the bucket → give the destination bucket name

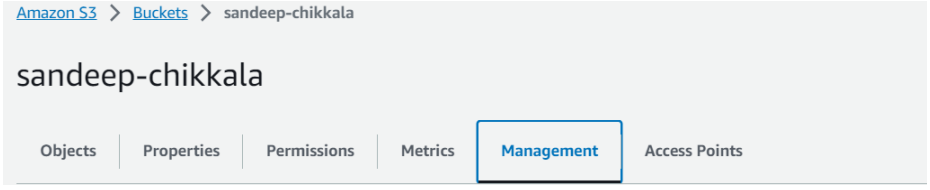

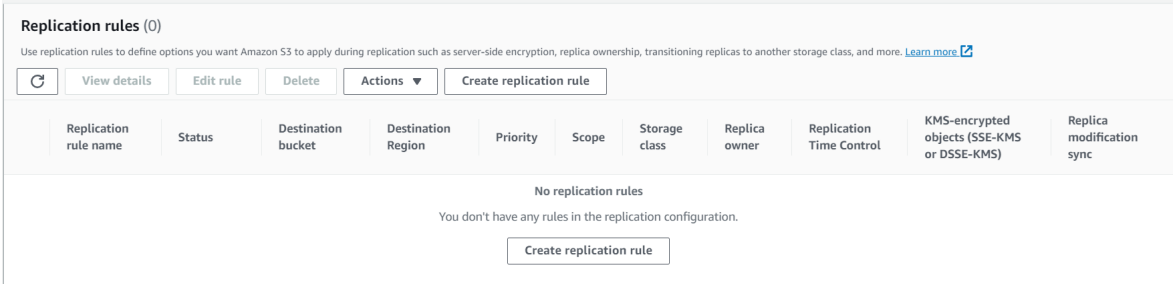

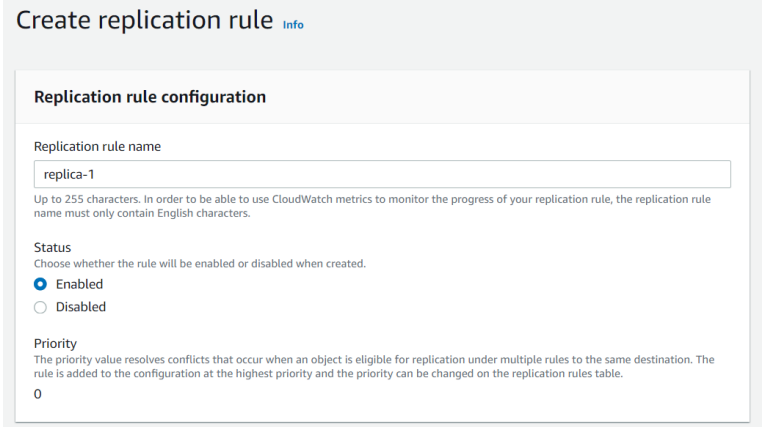

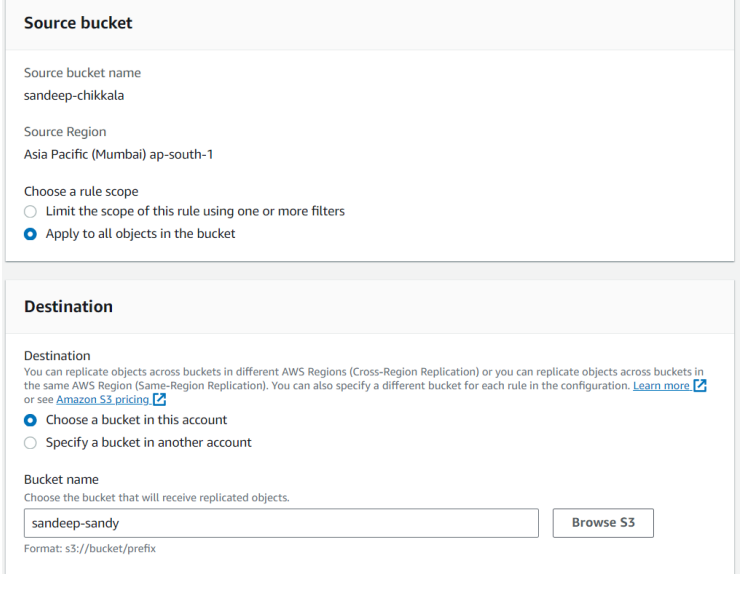

In IAM role → choose form existing IAM roles → select Create new role and save

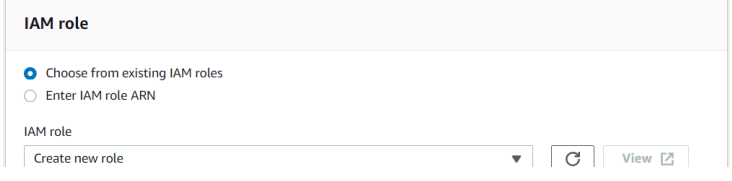

After perform above steps, click on save, it will ask yes (or) no, select yes

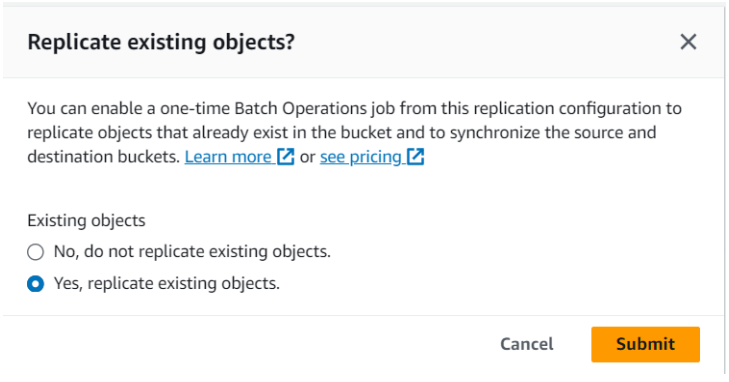

In the completion report, give bucket-2. For that click on browse S3, select the bucket and click on save

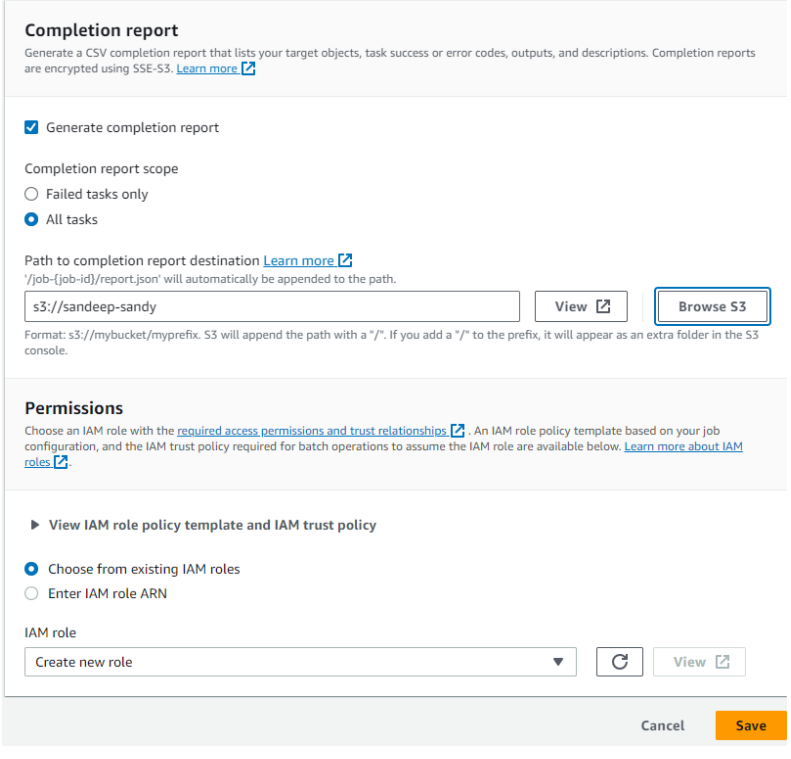

Step - 3 :

- So, whatever the step-2, we did for 1st bucket, now we have to do it 2nd bucket. Same like step-2
- So, if you upload a file, based on the size it will takes the time i.e. late replicas
- So, this CRR is worked in different accounts/different regions (or) same regions

Step - 4 :

Now, upload a file in 2nd bucket and check in 1st bucket and enable the show version. You will get the files

Here, overall we can upload the multiple files. But for deleting, it won't replicate. i.e. if you delete a file in

bucket-1 the changes didn't reflect in bucket-2

HOSTING THE STATIC WEBSITE MANUALLY IN S3

Websites are two types

- Static Websites
- Dynamic Websites

For hosting the Static websites we're using S3

- First, we need code
- Create a bucket, same like above one → go to properties → click on static website → edit

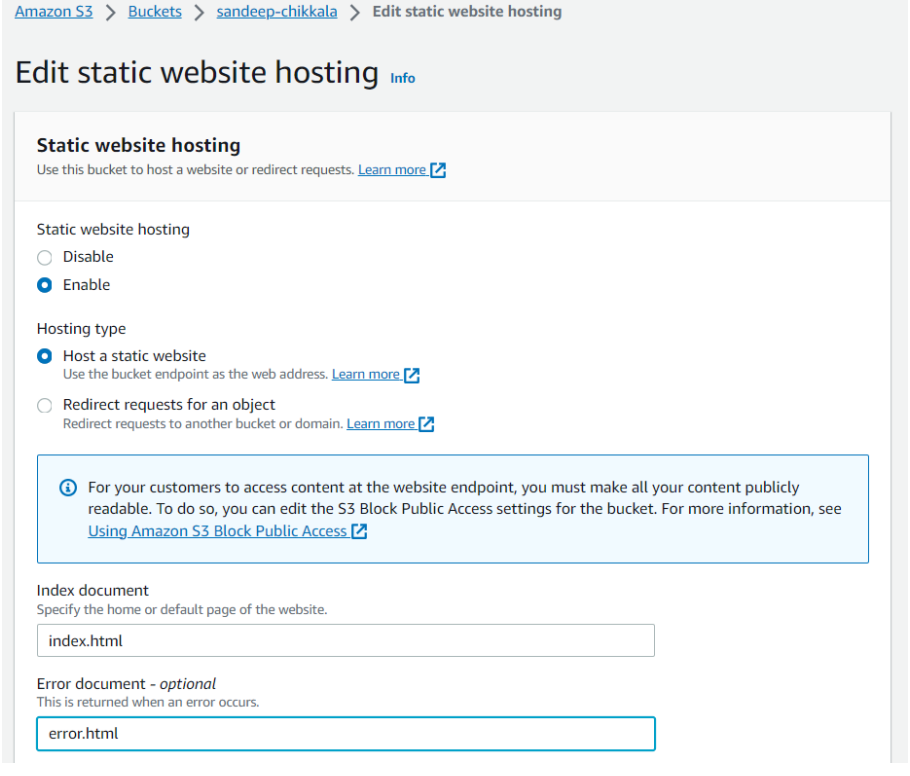

  - So, here for any website we are having home page. i.e. index.html
  - If the site is in under maintenance. So, that issue users want to know means that time they have to visible like a page called "error.html"
- After perform above data, click on save changes

So, first create 2 files in local with the same names

- index.html → <html><h1> hi this is static website </h1></html> → creating html page
- error.html → <html><h1> error </h1></html> → creating error page

Now, upload this 2 files in the bucket

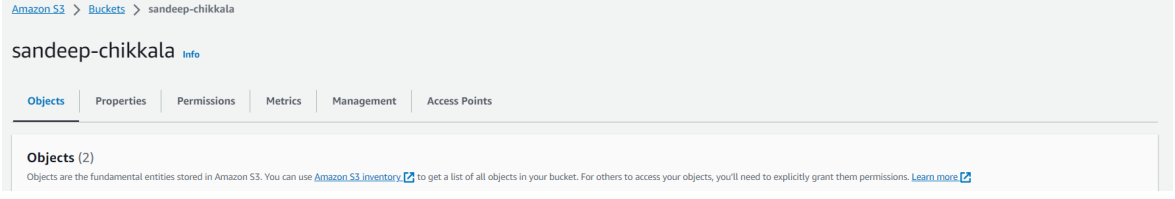

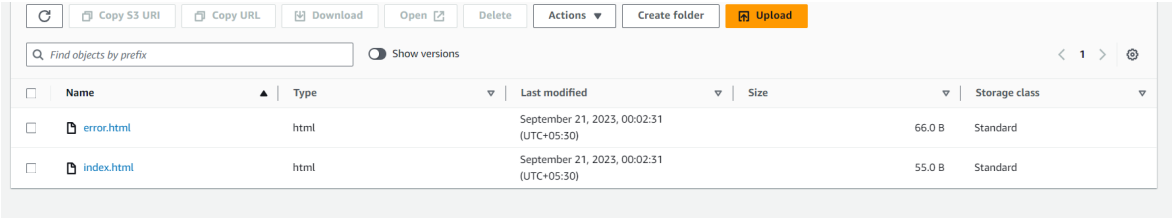

- So, when you select the object and copy the URL in browser, you will get output

(or)

- Go to properties → static hosting → and from we can copy the URL and paste in browser
- But, here we are getting permission denied errors.

For that, we have to do two things

1. Select the files inside objects and click on actions → select option "make public using ACL"

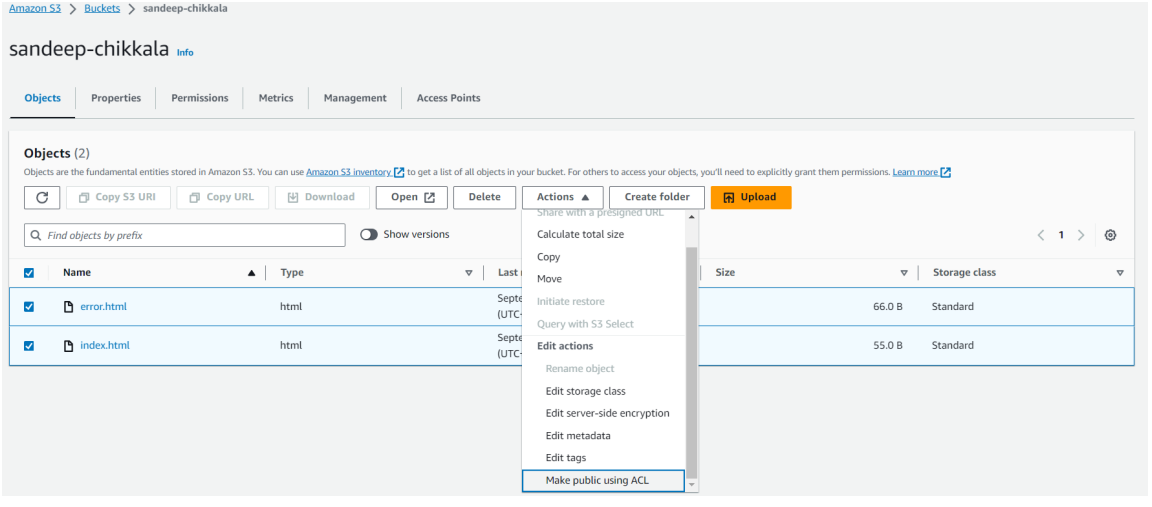

2. Go to Permissions → select block public access → edit

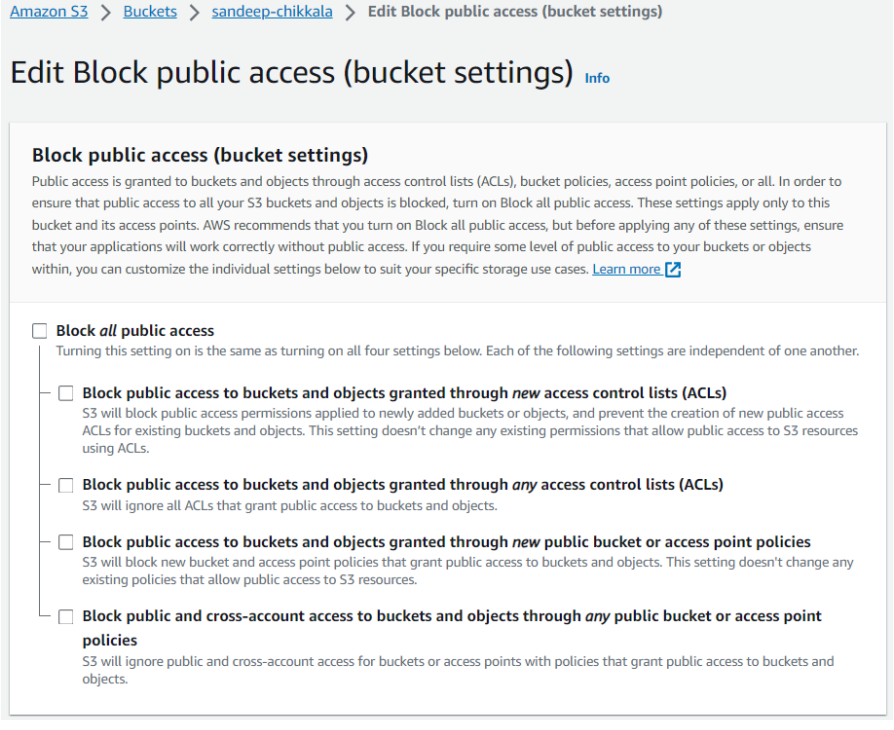

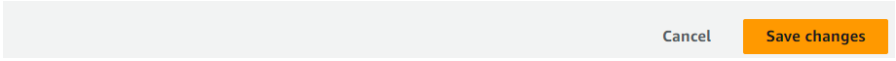

Here, uncheck everything and click on save changes

After doing changes, we will get like below image

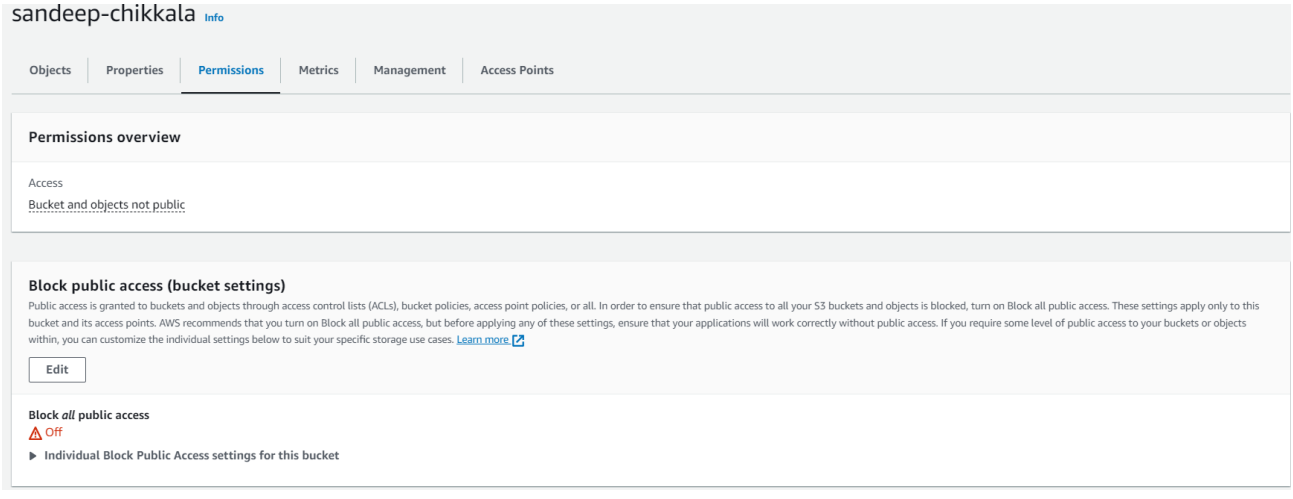

So, now you can access the static websites in browser to see the output

S3 - CLI

Through CLI, w can create, delete the buckets (or) objects

CLI Syntax : cloud service command

1. See list of buckets
    - aws s3 ls
2. Run - aws configure
    - Here, you have to provide access key and secret key
    - For that using IAM we have to create
    - Go to user → create user → After creation is completed → click on user → click on security credentials → click on access key → select CLI (or) third party service → create
    - Enter the details like this in the given below

```
[root@ip-172-31-6-121 ~]# aws configure
AWS Access Key ID [None]: AKIA3I5JELNFMM3RGVUP
AWS Secret Access Key [None]: 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
Default region name [None]: ap-south-1
Default output format [None]: table
```

1. Create a bucket
    - aws s3 mb s3://Bucketname
    - Here, mb → make bucket
    - Now, perform, aws s3 ls → you are having buckets
2. Delete the bucket
    - aws s3 rb s3://Bucketname
    - Here, rb → remove bucket
3. See the list of files/objects inside the bucket
    - aws s3 ls s3://Bucketname

4. Send a file from server to S3
    - touch aws.jpg
        - aws s3 cp aws.jpg s3://Bucketname
    - Check in S3, you will have the data
5. Download the bucket inside file in server
    - Go to S3 → select file → Copy S3 URI
        - aws s3 cp "S3URI" .
    - ll → you are having bucket files in your server
6. Delete the objects inside the bucket
        - aws s3 rm s3://Bucketname --recursive
    - Check in s3
7. Delete the bucket & file at a time
        - aws s3 rb s3://Bucketname --force
    - file deleted
    - Bucket deleted

Sync : It is a command used to send a files from one bucket to another bucket